

THE EXPERT'S VOICE® IN DB2

Beginning DB2

From Novice to Professional

*Everything you need to get up and running
with DB2 for Linux, UNIX, and Windows*

Grant Allen

Apress®

Beginning DB2

From Novice to Professional



Grant Allen

Beginning DB2: From Novice to Professional

Copyright © 2008 by Grant Allen

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN-13: 978-1-59059-942-6

ISBN-10: 1-59059-942-X

ISBN-13 (electronic): 978-1-4302-0548-7

ISBN-10 (electronic): 1-4302-0548-2

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Lead Editor: Jonathan Gennick

Technical Reviewer: Fred Sobotka

Editorial Board: Clay Andres, Steve Anglin, Ewan Buckingham, Tony Campbell, Gary Cornell, Jonathan Gennick, Matthew Moodie, Joseph Ottinger, Jeffrey Pepper, Frank Pohlmann, Ben Renow-Clarke, Dominic Shakeshaft, Matt Wade, Tom Welsh

Project Manager: Kylie Johnston

Copy Editor: Nancy Sixsmith, ConText Editorial Services, Inc.

Associate Production Director: Kari Brooks-Copony

Production Editor: Kelly Winkquist

Compositor: Diana Van Winkle

Proofreader: Linda Seifert

Indexer: Becky Hornyak

Artist: Diana Van Winkle

Cover Designer: Kurt Krames

Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail orders-ny@springer-sbm.com, or visit <http://www.springeronline.com>.

For information on translations, please contact Apress directly at 2855 Telegraph Avenue, Suite 600, Berkeley, CA 94705. Phone 510-549-5930, fax 510-549-5939, e-mail info@apress.com, or visit <http://www.apress.com>.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at <http://www.apress.com/info/bulksales>.

The information in this book is distributed on an “as is” basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

The source code for this book is available to readers at <http://www.apress.com>.

For Lindsay, who laughed at all the right times.

Contents at a Glance

About the Author	xix
About the Technical Reviewer	xxi
Acknowledgments	xxiii
Introduction	xxv

PART 1 ■ ■ ■ Getting Started

■ CHAPTER 1	What Is DB2 and Where Can I Get It?	3
■ CHAPTER 2	Installing DB2 for Linux and Windows	11

PART 2 ■ ■ ■ Beginning Administration with DB2 Express Edition

■ CHAPTER 3	DB2 Wizards: The Fast Path to Working with Your New Database ...	39
■ CHAPTER 4	Controlling Your Environment with DB2 Control Center and Command-Line Tools	63

PART 3 ■ ■ ■ Database Fundamentals with DB2 Express Edition

■ CHAPTER 5	IBM Data Studio	89
■ CHAPTER 6	SQL for DB2: Part 1	111
■ CHAPTER 7	SQL for DB2: Part 2	127
■ CHAPTER 8	Developing Functions and Stored Procedures in DB2	143
■ CHAPTER 9	Controlling Access to DB2	163
■ CHAPTER 10	Securing Data in DB2	177
■ CHAPTER 11	Using XML with DB2	195
■ CHAPTER 12	Indexes, Sequences, and Views	215

PART 4 ■ ■ ■ Programming with DB2 Express Edition

■ CHAPTER 13	PHP with DB2	241
■ CHAPTER 14	DB2 on Rails	263
■ CHAPTER 15	DB2 Development with Microsoft Visual Studio.NET	279
■ CHAPTER 16	Developing Java Applications with DB2	301
■ CHAPTER 17	Exploring More DB2 Development Options	323

PART 5 ■ ■ ■ Ongoing Database Administration with DB2

■ CHAPTER 18	Database Creation and Configuration	337
■ CHAPTER 19	DB2 Backup and Recovery	359
■ CHAPTER 20	Sharing Your DB2 Environment	385
■ CHAPTER 21	Moving Data in Bulk with DB2	407
■ CHAPTER 22	Working with Design	427
■ CHAPTER 23	Moving Up to More Advanced DB2	441
■ CHAPTER 24	Monitoring and Tuning DB2 Performance	459
■ INDEX	491

Contents

About the Author	xix
About the Technical Reviewer	xxi
Acknowledgments	xxiii
Introduction	xxv

PART 1 ■ ■ ■ Getting Started

■ CHAPTER 1	What Is DB2 and Where Can I Get It?	3
	Who Should Read This Book?	3
	Why Choose DB2?	4
	DB2 Has History on Its Side	4
	DB2 Is at the Forefront of Database Technology Today	4
	DB2 Will Be Even Bigger in the Future	4
	Choosing the Right DB2 Edition For You	5
	Obtaining DB2: A World of Choice	6
	Summary	9
■ CHAPTER 2	Installing DB2 for Linux and Windows	11
	Installing on Windows	11
	Installing on Linux	11
	Unpacking Installation Downloads	12
	Using the DB2 9 Discovery Kit DVD	12
	Checking Preinstallation Requirements	13
	Memory	13
	Disk	13
	Users and Groups	14
	Other Configuration Prerequisites	15
	Proceeding with the Installation	15
	Graphical Installation Under Windows and Linux	15
	Command-Line Installation Under Linux	32
	Uninstalling DB2	34
	Reviewing Your DB2 Installation	35
	Summary	35

PART 2 ■ ■ ■ Beginning Administration with DB2 Express Edition

■ CHAPTER 3	DB2 Wizards: The Fast Path to Working with Your New Database	39
	DB2 First Steps	39
	Launching DB2 First Steps Under Windows	40
	Launching DB2 First Steps Under Linux	41
	Working with DB2 First Steps	43
	Creating the SAMPLE Database	45
	The DB2 Control Center	47
	Starting the DB2 Control Center from First Steps	47
	Starting the DB2 Control Center in Other Ways	48
	Selecting Your Preferred DB2 Control Center View	49
	Browsing the DB2 Control Center Object Views	50
	Invoking the Create Table Control Center Wizard	52
	Using the DB2 Command Editor	55
	Command-Line Control of Your DB2 Environment	58
	Powering Up the Command Line Processor	58
	Learning About the Commands	59
	Connecting and Disconnecting	60
	Summary	61
■ CHAPTER 4	Controlling Your Environment with DB2 Control Center and Command-Line Tools	63
	Getting Started as the Right User	63
	Starting and Stopping DB2 Processes	64
	From the Control Center	64
	From the Command Line	66
	Server Components	70
	DB2 Instances and the DAS Instance	70
	The DB2 Governor	71
	The DB2 License Server	71
	The DB2 Management Service	71
	The DB2 Remote Command Server	72
	The DB2 Security Server	72
	The DB2 Fault Monitor Coordinator	72

Configuring and Changing Your DB2 Instances and Databases	73
Two Different Levels of Configuration	74
Configuring Your DB2 Instances from the Control Center	75
Configuring a DB2 Database from the Control Center	78
Configuring and Changing Your DB2 Instances from the CLP	79
Configuring and Changing a DB2 Database from the CLP	81
More GUI Tools to Help You Manage DB2	82
More Tools Available from the Command Line	82
The Version Information Utility: DB2LEVEL	82
The DB2 Registry Utility: DB2SET	83
The Instance Listing Utilities: DB2ILIST and DASLIST	83
The DB2 Problem Determination Tool: DB2PD	84
Summary	85

PART 3 ■ ■ ■ Database Fundamentals with DB2 Express Edition

■ CHAPTER 5	IBM Data Studio	89
	Sourcing the Data Studio	90
	Installation Highlights	91
	Multiple Installers for Data Studio	91
	The Power of Eclipse-Based IDE Installations	92
	Running Data Studio	94
	Starting a New Data Project	95
	New Project Connection Management	97
	Adding Objects to Your Project	100
	New SQL Statement Wizard	100
	New Stored Procedure Wizard	104
	Summary	109
■ CHAPTER 6	SQL for DB2: Part 1	111
	Manipulating Data with SQL	111
	Using Select Statements	112
	Using Insert Statements	124
	Using Update Statements	125
	Using Delete Statements	126
	Summary	126

CHAPTER 7	SQL for DB2: Part 2	127
	Creating and Managing DB2 Tables	127
	Data Types in DB2	127
	Your First Table	130
	Naming DB2 Tables	131
	Specifying Column Details and Constraints	132
	Working with Null Data	132
	Primary Key Constraints	133
	Referential Constraints	134
	Check Constraints	136
	Disabling Constraints and Constraint Deferral	137
	Automatic Value Generation for Columns	138
	Specifying Table Storage Characteristics	140
	Other Table Characteristics	141
	Other Techniques for Table Creation	142
	Summary	142
CHAPTER 8	Developing Functions and Stored Procedures in DB2	143
	DB2 Stored Procedures	143
	Creating Stored Procedures in DB2	143
	Dropping Procedures	152
	Managing Procedures	154
	Further Reading on DB2 Stored Procedures	155
	User-Defined Functions in DB2	156
	Types of User-Defined Functions in DB2	156
	Creating User-Defined Functions in DB2	157
	Calling User-Defined Functions in DB2	159
	Managing and Deleting User-Defined Functions	161
	Summary	161
CHAPTER 9	Controlling Access to DB2	163
	DB2 Building Blocks for Authentication	164
	Configuring Where and When DB2 Authenticates Users	164
	Choosing Your Authentication Option	167
	DB2 Instance-Level Privileged Groups	168

Changing Your DB2 Authentication Parameters	169
Using Control Center to Manage Authentication Parameters	170
Using DB2 CLP to Manage Authentication Parameters	172
Group Authentication Issues	173
Configuring DB2 to Use Local or Global Groups	173
64 Groups Limitation	174
Support for Windows AD Domain Features	174
Security Context of the DB2 Server	174
Summary	175
 CHAPTER 10 Securing Data in DB2	177
Managing Authorities and Privileges in DB2	177
Preparing to Use Authorities and Privileges	177
Database-Level Authorities	177
Group Behavior in DB2	182
Roles in DB2	184
Managing Object Privileges in DB2	185
Working with Privileges on DB2 Objects	185
Granting Object Privileges to Users	187
Label-Based Access Control in DB2	190
LBAC in Action	192
Summary	194
 CHAPTER 11 Using XML with DB2	195
Exploring XML in the Sample Database	195
Querying Your XML Data	198
Using XQuery for XML	199
Using XPath Queries for XML	204
More pureXML Features for Querying Data	207
Changing XML Data	208
Inserting XML Data	208
XML Schema Registration in DB2	210
Updating XML Data	212
Deleting XML Data	213
Summary	214

■ CHAPTER 12	Indexes, Sequences, and Views	215
	Working with Indexes	215
	Creating Indexes	215
	Enforcing Unique Values	217
	The Important Case of Foreign Keys	218
	Understanding Other Index Features	219
	Using Design Advisor	221
	Index Wrap Up	227
	Working with Sequences	228
	Creating Sequences	228
	Altering Sequences	230
	Using Sequences	231
	Sequence Wrap Up	233
	Working with Views	233
	Defining Views	233
	Further View Options	235
	View Wrap Up	237
	Summary	237

PART 4 ■ ■ ■ Programming with DB2 Express Edition

■ CHAPTER 13	PHP with DB2	241
	PHP and Zend Core for DB2	241
	Starting an Install of Zend Core	242
	Continuing the Installation	244
	Post-Installation Checks	247
	Developing with the IBM_DB2 PHP Extensions	248
	Connecting to a Database	248
	Managing IBM_DB2 PHP Connections	252
	Working with Ad Hoc SQL Statements	252
	Working with Prepared Statements and Stored Procedures	255
	Other IBM_DB2 PHP Functions	257
	Developing with the Zend Framework	258
	Configuring the Zend Framework	258
	Managing Connections with the Zend Framework	259

Executing Statements Using the Zend Core Framework	260
Other Major Classes in the Zend Core Framework	262
Summary	262

■ CHAPTER 14 DB2 on Rails 263

Installing Ruby and Rails	263
Developing with Ruby, Rails, and DB2	267
Creating a Project with Rails	267
Connecting a Rails Project to DB2	268
Designing DB2 Tables in Rails	269
Executing a Migration in Ruby for DB2	271
Evolving Your DB2 Schema with Ruby on Rails	272
The Schema_Info Table for Ruby on Rails	273
Reverting to Earlier Schema Incarnations	274
Ruby on Rails Scaffolding for DB2	274
Generating the Scaffolding	275
Ruby on Rails Scaffolding in Action with DB2	275
Summary	277

■ CHAPTER 15 DB2 Development with Microsoft Visual Studio.NET 279

Installing IBM Database Add-Ins	279
Registering the IBM Data Server Provider for .NET	281
Testing DB2 .NET Connectivity	282
Completing the Installation	283
Testing the Database Add-Ins	283
Managing DB2 Objects	286
Creating a New DB2 View Object	287
Building the DB2 Database Project	289
Building .NET Windows Applications with DB2	290
Binding DB2 Objects to Windows Controls	291
Specifying a DB2 Connection for a Windows C# Project	292
Exploring a Connection's Related Objects	294
Building and Running the Windows Project	294
Building .NET Web Applications with DB2	295
Binding DB2 Objects to Web Controls	296
Building and Running the Web Project	298
Summary	299

CHAPTER 16	Developing Java Applications with DB2	301
	Writing JDBC Applications for DB2	302
	JDBC Driver Types	302
	Writing the First Java DB2 Application	303
	Refactoring the Connection Code	305
	Working with Data	307
	Working with DB2 JDBC Features	309
	Developing Java Stored Procedures	311
	Writing the Stored Procedure Java Class	311
	Deploying the Java Stored Procedure Class	312
	Writing the Stored Procedure	313
	Testing the Java DB2 Stored Procedure	314
	Getting Started with DB2 pureQuery for Java	315
	Creating a pureQuery-Enabled Project	315
	Autogenerating pureQuery Java Mapping Classes	318
	Examining pureQuery Results	321
	Summary	322
 CHAPTER 17	 Exploring More DB2 Development Options	 323
	Perl	323
	Installing the Perl DBI and DBD::DB2 Modules	324
	Getting Connected	327
	Executing SQL Statements	328
	Going Further with Perl and DB2	330
	Python	330
	Installing the Python DB2 Driver	330
	Getting Connected	332
	Executing SQL Statements	332
	Going Further with Python and DB2	332
	Other Languages	333
	C and C++	333
	COBOL	333
	Rexx	333
	Visual Basic and Visual Basic.NET	334
	Other Samples	334
	Summary	334

PART 5 ■ ■ ■ Ongoing Database Administration with DB2

■ CHAPTER 18	Database Creation and Configuration	337
	Building Blocks of DB2 Databases	337
	Creating DB2 Databases	338
	Specifying an Alias	339
	Handling Languages and Text in DB2 Databases	340
	Managing Storage for the DB2 Database	342
	Understanding Tablespace Use in DB2	342
	Choosing the Right Tablespace Storage Management Type	343
	Using System Managed Storage Tablespaces	344
	Using Database Managed Storage Tablespaces	345
	Tuning Initial Tablespace Parameters	346
	Introducing Buffer Pools	347
	Moving Beyond Database, Tablespace, and Buffer Pool Basics	348
	Creating Databases with Many Options	349
	Altering Databases	349
	Working with Tablespaces and Buffer Pools	350
	Dropping Databases	356
	Summary	357
■ CHAPTER 19	DB2 Backup and Recovery	359
	Why You Should Care	359
	Permissions for Backup and Recovery	359
	Understanding DB2 Logging	360
	General Logging Principles	360
	Log File Types	361
	Circular Logging	361
	Archive Logging	362
	Infinite Logging	365
	Log Housekeeping	366
	DB2 Backups	366
	Anatomy of a Backup	367
	Taking Backups from the Control Center	368
	Advanced Backup Options from the Command Line	372

Database Recovery with DB2	375
Database Recovery for Free	375
Restoring a Database from Backup	375
Roll Forward Recovery from a Backup	376
Database Recovery Using the Control Center	377
Advanced Recovery Options from the Command Line	381
Summary	383
 CHAPTER 20 Sharing Your DB2 Environment	385
DB2 Clients on Linux and Windows	385
Using the DB2 Express-C Edition Installer	386
Using the Dedicated DB2 Client Installers	387
Deciding Which Components to Install	393
Configuring the DB2 Client for Your Database	394
The DB2 Client Configuration Assistant	394
Command-Line DB2 Client Configuration	400
Other Connectivity Options	405
DB2 Runtime Client	405
Type 4 JDBC Driver	405
Third-Party Connectivity Options	405
Summary	406
 CHAPTER 21 Moving Data in Bulk with DB2	407
File Formats for Moving Data	407
ASCII Format	407
Delimited Format	407
Lotus 1-2-3 Worksheet Format	408
Integration Exchange Format	408
Exporting Data	408
Performing a Simple Export	409
Exporting LOBs	409
Exporting XML Objects	411
Exporting via the Control Center	412
Importing Data	414
Performing a Simple Import	414
Dealing with Import Warnings and Errors	416
Importing a Mix of New and Updated Data	418
Performing More-Complex Imports	419
Graphical Import Using the Control Center	421

Loading Data with the Load Utility	423
Other Data-Movement Tools	426
Summary	426
CHAPTER 22 Working with Design	427
Database Design Overview	428
Statement of Requirements	429
Identification of Users and Use Cases	430
Analysis of Use Cases and Data Requirements	432
Finding Relationships in the Data	433
Building the Logical Data Model	435
Translating a Logical Model to a Physical Model	436
Building the Physical Model in DB2	436
Model Review	440
Summary	440
CHAPTER 23 Moving Up to More Advanced DB2	441
Using the db2look and db2move Utilities	441
Using db2look	441
Using db2move	443
Using db2look from the Control Center	445
DB2 Replication	446
Preparing for Replication	446
Configuring Replication via the Replication Center	447
Testing Replication in Action	452
Data Row Compression	454
Data Row Compression Design	455
Estimating Compression Savings	456
Compressing a Table	457
Limitations	458
Summary	458
CHAPTER 24 Monitoring and Tuning DB2 Performance	459
Starting with a Sensible Approach to Tuning	459
Proactive Tuning with the Configuration Advisor	460
Understanding Isolation Levels in DB2	464
Understanding Locking in DB2	464
Using DB2's Optimistic Locking Features	466
Continuing with the Configuration Advisor	466

- Understanding Self-Tuning Memory in DB2 468
 - Traditional DB2 Memory Management Model 468
 - Activating STMM Mode 469
 - Self-Tuning Memory Management Mechanics 470
 - Using the Memory Visualizer 471
- Monitoring and Tuning Database Activity 472
 - Activating Monitor Switches 473
 - Using DB2 Event Monitors 473
 - Using DB2 Snapshot Monitors 478
- Tuning the DB2 Optimizer with RUNSTATS 481
 - Automated RUNSTATS in DB2 481
 - Manually Running the RUNSTATS Command 484
- Using the REORGCHK and REORG Utilities 485
 - Working with REORGCHK 485
 - Working with REORG 487
- Summary 489

- INDEX 491

About the Author

■ **GRANT ALLEN** has worked in the IT field for nearly 20 years, most recently as Chief Technology Officer for a leading Australian software vendor, before taking on his current role at Google. His work has involved private enterprise, academia, and government around the world—consulting on large-scale systems design, development, performance, data warehousing, content management, and collaboration. Grant is a frequent speaker at conferences and industry events on topics such as **data mining**, compliance, relational databases, collaboration technologies, the business of technology, and more. He is now a team leader at Google, using database technologies to tackle problems of Google-scale and beyond.

6.6.12